



Software

# BIGDL: DISTRIBUTED DEEP LEARNING ON APACHE SPARK

Yiheng Wang (yiheng.wang@intel.com)

Big Data Technologies, Software and Service Group, Intel

# Introduction

## Intel Big Data Technology team

- Active open source development
- Spark, Hadoop, HBase, Hive, Sentry, Storm, etc
- ~30 project committers in the team

## My focusing area

- Large scale machine learning, deep learning
- Next generations of Big Data analytics solutions with Intel customers

# BigDL

A distributed deep learning framework on Apache Spark

<http://www.github.com/intel-analytics/BigDL>

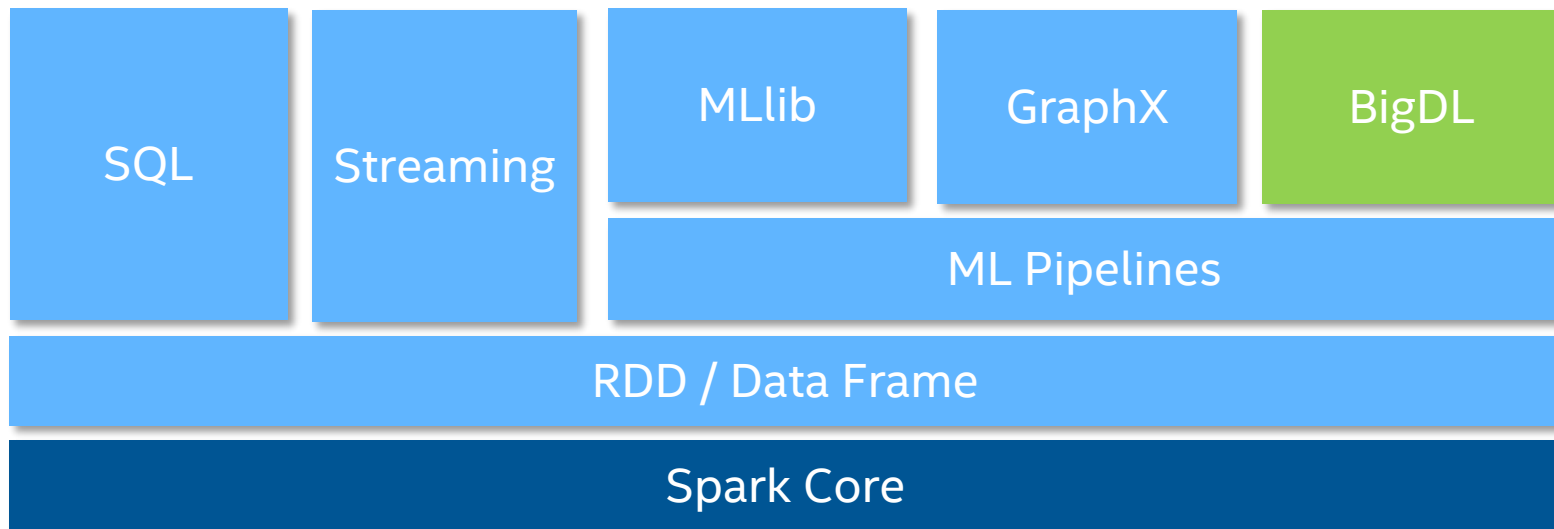
# Outline

- Overview
- Install and Run BigDL
- Define models
- Train and evaluate models
- Model tuning

# OVERVIEW

An overview of BigDL

# Build on Apache Spark



# There're a lot of deep learning solutions



Caffe

***mxnet***

theano



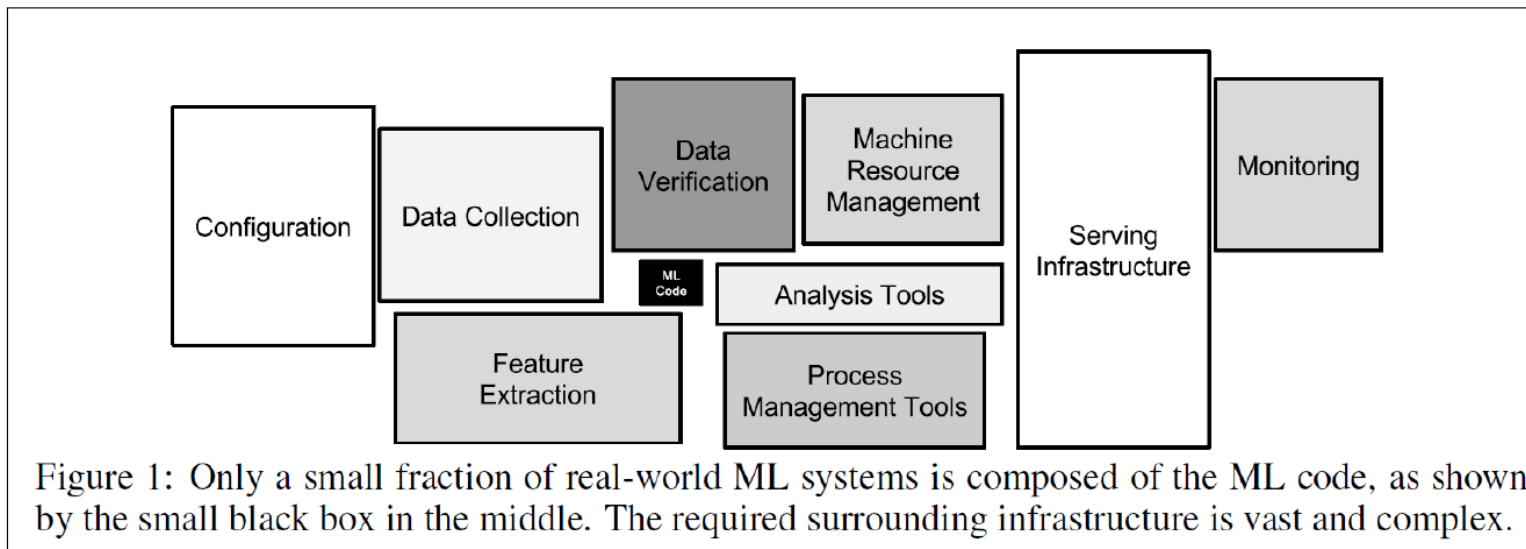
CNTK

# BigDL

A scalable and easy solution for  
deep learning on Big Data



# Build an End-2-end Solution



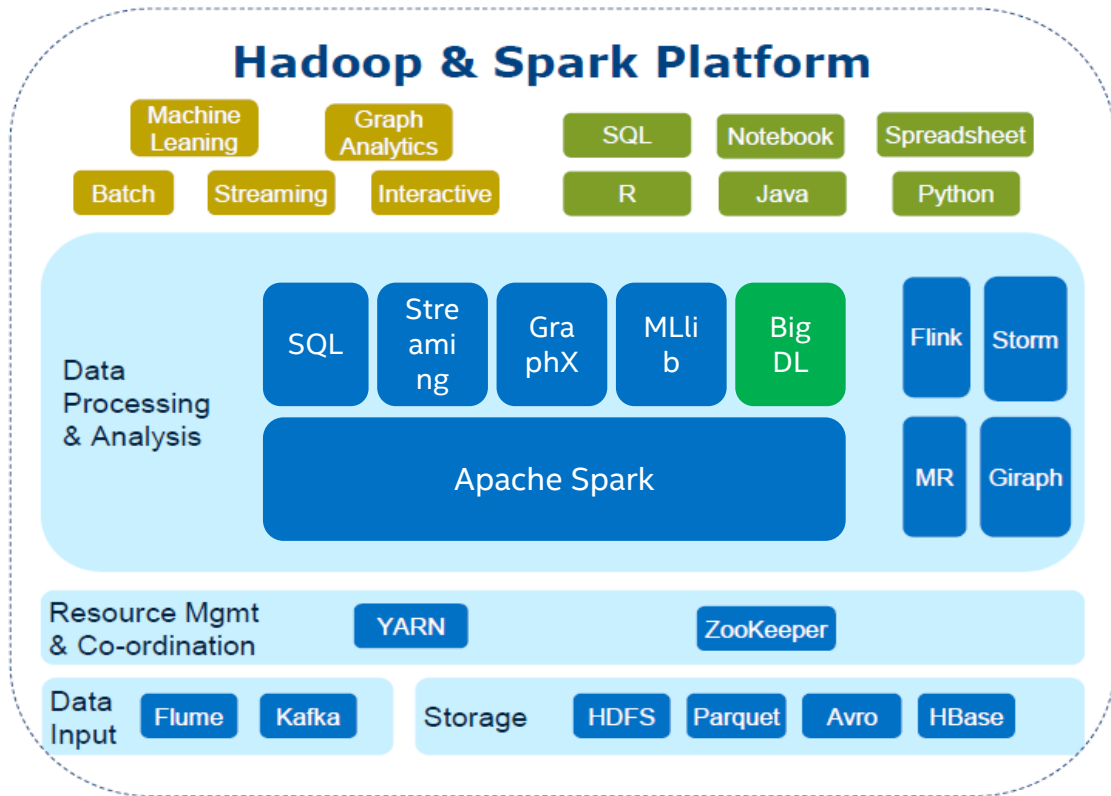
“Hidden Technical Debt in Machine Learning Systems”,  
Google, NIPS 2015 Paper

# Build an End-2-end Solution

## Practical challenges:

- compatible with different data source
- performance and scalability
- stability & fault tolerant
- data management / pre-processing
- resource sharing
- programming tools / languages
- ...

# Build an End-2-end Scalable Solution






# BigDL is easy to use

- A friendly API compatible with Torch
- Provide Scala and Python programming API

# BigDL is easy to deploy

## Real out-of-box



# BigDL is easy to deploy

Public cloud blogs (See <https://github.com/intel-analytics/BigDL/wiki/powered-by>)

- Intel's BigDL on Databricks
- Use BigDL on AZure HDInsight
- BigDL on AliCloud E-MapReduce (in Chinese)
- Running BigDL, Deep Learning for Apache Spark, on AWS
- Running BigDL on Microsoft Data Science Virtual Machine
- Using Apache Spark with Intel BigDL on Mesosphere DC/OS by Lightbend

People use BigDL to build applications

# Rich deep learning feature support

## Layers

- More than 100 (Linear, Conv2D, Conv3D, Embedding, Recurrent)

## Loss function

- Dozens of loss functions

## Optimization algorithm

- SGD, Adagrad, Adam, Adamax, RMSProp, Adadelata

## Distributed Training / Inference

## Save and Load model files

- Also include torch / caffe / tensorflow

# High performance from your server

- Powered by Intel Math Kernel Library
- Extremely high performance on Xeon CPUs
  - Order of magnitude faster than out of box caffe / torch / tensorflow
  - Comparable with GPU (same generation)
- Good scalability
  - Hundreds of nodes



# INSTALL AND RUN BIGDL

How to install and run bigdl on your cluster

# Get executable BigDL

- Download
- Maven / Sbt
- Pip install
- Build yourself

# Download

- Download Page (<https://github.com/intel-analytics/BigDL/wiki/Downloads>)
  - Linux x64 and Mac OS
  - Windows(WIP)
- Stable release and nightly build
- Python development / Run examples

# Maven / SBT

- <https://github.com/intel-analytics/BigDL/wiki/Build-Page#linking>
- Snapshot, release
- Java/Scala development

```
<dependencies>
  <dependency>
    <group>com.intel.analytics.bigdl</group>
    <artifactId>bigdl-SPARK_(1.5/1.6/2.0/2.1)</artifactId>
    <version>0.1.1</version>
  </dependency>
</dependencies>
```

# Pip Install

See <https://github.com/intel-analytics/BigDL/wiki/Install-BigDL-via-pip>

- Download Spark2.x

```
wget https://d3kbcqa49mib13.cloudfront.net/spark-2.1.0-bin-hadoop2.7.tgz
```

- Extract the tar ball and set SPARK\_HOME

```
tar -zxvf spark-2.1.0-bin-hadoop2.7.tgz  
export SPARK_HOME=path to spark-2.1.0-bin-hadoop2.7
```

- Install BigDL 0.1.1 release via pip (we tested this on pip 9.0.1)

```
pip install --upgrade pip  
pip install BigDL==0.1.1rc0 # for Python 2.7  
pip3 install BigDL==0.1.1rc0 # for Python 3.n
```

- Launch with Python REPL / Jupyter

# Build yourself

- Customized configuration, e.g. JDK 8, Spark version
- Develop BigDL
- No need to pre-install MKL (MKL jar will be downloaded)

```
$ git clone https://github.com/intel-analytics/BigDL.git
$ cd BigDL

$ ./make-dist.sh # For Spark 1.5/1.6, Linux x64

$ ./make-dist.sh -P mac # For Spark 1.5/1.6, MacOS

$ ./make-dist.sh -P spark_2.x # For Spark 2.0/2.1, Linux x64

$ ./make-dist.sh -P mac -P spark_2.x # For Spark 2.0/2.1, MacOS
```

# Start your BigDL program

## Run scala code


```
spark-submit \  
  --master xxx  
  --jars path_to_big_dl_jar  
  --class main_class_full_name  
  --.....  
  your_project_jar  
  .....
```

## Run python code

```
spark-submit \  
  --master xxx  
  --jars path_to_big_dl_jar  
  --py-files path_to_big_dl_python_zip  
  your_python_file  
  .....
```

In BigDL 0.1.0 and 0.1.1, you need to run **source bigdl.sh** before you run the spark-submit command


# How BigDL run on Apache Spark\*





# Sign up for free compute for BigDL

<https://software.intel.com/en-us/ai/frameworks/bigdl/remote-access>



Preregister for Free Compute for BigDL, sponsored by Intel, and provide feedback to help make BigDL better for new users. You don't need to share your code. Preference goes to those who share their BigDL story.

# DEFINE MODELS

How to define model in BigDL

## Input



# BigDL provides two kind of model definition style

- Sequential API
  - In sequential API, user add layer into some containers to build the model
- Functional API
  - In functional API, the model is described as a graph



# Define a model (Linear)

## Scala

```
val model = Sequential()  
model.add(Linear(4, 3))
```

## Python

```
model = Sequential()  
model.add(Linear(4, 3))
```



Simple Linear classification

$$Y = X * W + B$$




# Add activation functions

## Scala

```
val model = Sequential()  
model.add(Linear(4, 3))  
model.add(Sigmoid())
```

## Python

```
model = Sequential()  
model.add(Linear(4, 3))  
model.add(Sigmoid())
```




# Logistic Regression

## Scala

```
val model = Sequential()  
model.add(Linear(4, 1))  
model.add(Sigmoid())  
Model.add(Softmax())
```

## Python

```
model = Sequential()  
model.add(Linear(4, 1))  
model.add(Sigmoid())  
Model.add(Softmax())
```





# Another way to define Logistic Regression

## Scala

```
val linear = Linear(4, 1).inputs()
val sigmoid = Sigmoid().inputs(linear)
val softmax = Softmax().inputs(sigmoid)
val model = Graph(Seq[linear], Seq[softmax])
```

## Python

```
linear = Linear(4, 1)()
sigmoid = Sigmoid()(linear)
softmax = Softmax()(sigmoid)
model = Model([linear], [softmax])
```





# Define a model with branches

## Sequential

```
branch1 = Sequential().add(Linear(...)).add(ReLU())
```

```
branch2 = Sequential().add(Linear(...)).add(ReLU())
```


```
branches =  
ConcatTable().add(branch1).add(branch2)
```

```
val model = Sequential()
```

```
model.add(Linear(...))
```

```
model.add(ReLU())
```


```
model.add(branches)
```



# Define a model with branches

## Functional

```
linear1 = Linear(...>()  
relu1 = ReLU()(linear1)  
linear2 = Linear(...)(relu1)  
relu2 = ReLU()(linear2)  
linear3 = Linear(...)(relu1)  
relu3 = ReLU()(linear3)  
model = Model(Seq[linear1], Seq[relu2, relu3])
```



# Define a model with merged branch

## Sequential

```
branch1 = Sequential().add(Linear(...)).add(ReLU())
```

```
branch2 = Sequential().add(Linear(...)).add(ReLU())
```

```
branches =  
ConcatTable().add(branch1).add(branch2)
```


```
val model = Sequential()
```

```
model.add(Linear(...))
```

```
model.add(ReLU())
```

```
model.add(branches)
```


```
model.add(CAddTable())
```



# Define a model with merged branch

## Functional

```
linear1 = Linear(...>()  
relu1 = ReLU()(linear1)  
linear2 = Linear(...)(relu1)  
relu2 = ReLU()(linear2)  
linear3 = Linear(...)(relu1)  
relu3 = ReLU()(linear3)  
add = CAddTable()(relu2, relu3)  
model = Model(Seq[linear1], Seq[add])
```



# Define a model with multiple inputs

## Sequential


```
model = Sequential()

branches = ParallelTable()

branch1 = Sequential().add(Linear(...)).add(ReLU())
branch2 = Sequential().add(Linear(...)).add(ReLU())

branches.add(branch1).add(branch2)


model.add(branches).add(CAddTable)
```



# Define a model with multiple inputs

## Functional

```
linear1 = Linear(...>()  
relu1 = ReLU()(linear1)  
linear2 = Linear(...>()  
relu2 = ReLU()(linear2)  
add = CAddTable()(relu1, relu2)  
model = Model(Seq[linear1, linear2], Seq[add])
```



# Model definition

**Let's take a look at some other layers**

# Convolution neural networks

## Convolution Layers

- Widely used in image related models (not limited)

1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved  
Feature



\*

1	0	-1
2	0	-2
1	0	-1



Images are from: <https://leonardoaraujosantos.gitbooks.io/artificial-intelligence/content/convolution.html>

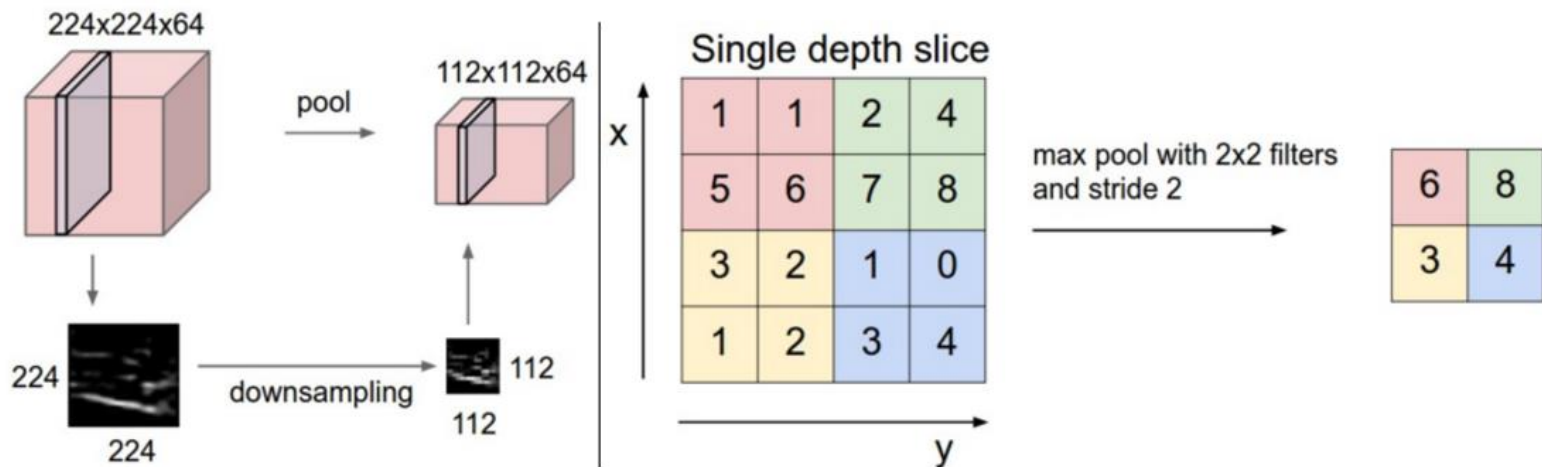


# Convolution neural networks

```
SpatialConvolution(  
    nInputPlane, nOutputPlane,  
    kernelW, kernelH,  
    strideW=1, strideH=1,  
    padW=0, padH=0,  
    nGroup=1,  
    wRegularizer=null, bRegularizer=null,  
    initWeight=null, initBias=null, initGradWeight=null, initGradBias=null  
)
```

# Convolution neural networks

## Pooling




The image is from: [https://leonardaraujosantos.gitbooks.io/artificial-intelligence/content/pooling\\_layer.html](https://leonardaraujosantos.gitbooks.io/artificial-intelligence/content/pooling_layer.html)

# Convolution neural networks

```
SpatialMaxPooling(  
    kW, kH,  
    dW=1, dH=1,  
    padW=0, padH=0,  
    ceilMode=false  
)
```

# RNN




The repeating module in a standard RNN contains a single layer.

Source: <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

# RNN

```
model.add(  
    Recurrent[Float]()  
        .add(  
            RnnCell[Float](inputSize, outptuSize, Tanh[Float]()  
        )  
    )  
)
```

# LSTM




The repeating module in an LSTM contains four interacting layers.

Source: <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

# LSTM

```
model.add(  
    Recurrent[Float]()  
        .add(  
            LSTM[Float](inputSize, outptuSize)  
        )  
)
```

# Embedding Layer





# Embedding Layer

LookupTable(

  nIndex: Int, nOutput: Int,

  paddingValue: Double = 0,

  maxNorm: Double = Double.MaxValue,

  normType: Double = 2.0,

  shouldScaleGradByFreq: Boolean = false,

  wRegularizer: Regularizer[T] = null

)

A close-up photograph of a calico cat lying on its back on a light-colored surface. The cat's front paws are held up near its face, and its hind legs are also extended upwards. The cat has a white base coat with patches of black and orange. Its eyes are wide open and looking towards the camera. The text "TAKE A BREAK" is overlaid in large, bold, blue capital letters on the left side of the image.

**TAKE A BREAK**

# TRAIN AND EVALUATE MODEL

How to train a model and how to evaluate it

# We will take a look at

- How to prepare your data
- Define a training process
- Predict with your model

# Data preprocess

The raw data(image, audio, text) can not be used with model directly

- They need to be convert to tensors
- Preprocessing is often necessary
  - Normalization
  - Embedding
  - Scale
  - Crop
  - Augmentation

# Data preprocess

In Python, thanks to the rich data analytics libraries, you can do it easily

- Numpy, Pandas...

In Scala, BigDL provide several utilities to do preprocessing

```
trait Transformer[A, B] extends Serializable {  
  def apply(prev: Iterator[A]): Iterator[B]  
}
```

# Data preprocess in Scala

```
class PathToImage extends Transformer[Path, Image]
class ImageToArray extends Transformer[Image, Array]
class Normalizer extends Transformer[Array, Array]
class Cropper extends Transformer[Array, Array]
```

```
PathToImage -> ImageToArray -> Normalizer -> Cropper
```

```
val rddA : RDD[A] = ...
val tran : Transformer[A, B] = ...
val rddB : RDD[B] = rdd.mapPartitions(tran(_))
```

# Tensor

## Numpy NDarray for Python


```
np.array(  
    [  
        [1.0, 1.0, 1.0, 1.0]  
        [3.0, 3.0, 3.0, 3.0]  
    ]  
)
```

## Tensor for Scala

```
Tensor[Float](  
    T(  
        T(1.0f, 1.0f, 1.0f, 1.0f),  
        T(3.0f, 3.0f, 3.0f, 3.0f)  
    )  
)
```



# Sample



In distributed training or inference

- `RDD[Sample]`

# Let's prepare some data

## MNIST Dataset

- <http://yann.lecun.com/exdb/mnist/>

# THE MNIST DATABASE

## of handwritten digits

Yann LeCun, Courant Institute, NYU

Corinna Cortes, Google Labs, New York

Christopher J.C. Burges, Microsoft Research, Redmond

# Sandbox environments

<https://github.com/yiheng/OReillyAIConf#sandbox-environment>

# Take a look at MNIST data

```
%pylab inline
from bigdl.dataset import mnist

mnist_path = "datasets/mnist"
(train_images, train_labels) = mnist.read_data_sets(mnist_path, "train")
(test_images, test_labels) = mnist.read_data_sets(mnist_path, "test")

print train_images.shape
print train_labels.shape
print test_images.shape
print test_labels.shape

imshow(np.column_stack(train_images[0:10].reshape(10, 28,28)),cmap='gray'); axis('off')
print "groud true labels: "
print train_labels[0:10]
```

# Convert MNIST to RDD (code part 1)

```
from bigdl.util.common import Sample
from bigdl.dataset import mnist

def get_mnist(sc, mnist_path):
    # target is start from 0,
    (train_images, train_labels) = mnist.read_data_sets(mnist_path, "train")
    (test_images, test_labels) = mnist.read_data_sets(mnist_path, "test")
    training_mean = np.mean(train_images)
    training_std = np.std(train_images)
    rdd_train_images = sc.parallelize(train_images)
    rdd_train_labels = sc.parallelize(train_labels)
    rdd_test_images = sc.parallelize(test_images)
    rdd_test_labels = sc.parallelize(test_labels)
```

## Convert MNIST to RDD (code part 2)


```
    rdd_train_sample = rdd_train_images.zip(rdd_train_labels).map(lambda (features,
label):
    Sample.from_ndarray((features - training_mean)/training_std, label + 1))
    rdd_test_sample = rdd_test_images.zip(rdd_test_labels).map(lambda (features, label):
    Sample.from_ndarray((features - training_mean)/training_std, label + 1))
    return (rdd_train_sample, rdd_test_sample)

(train_data, test_data) = get_mnist(sc, mnist_path)
print train_data.count()
print test_data.count()
```

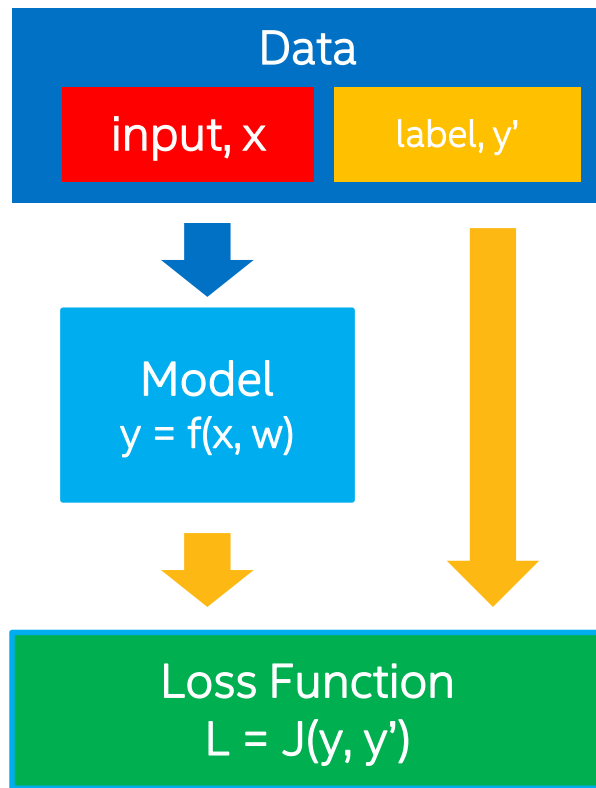

Define a training process

**Take a look at the theory first**

# How to train the model



<https://www.quora.com/Whats-the-difference-between-gradient-descent-and-stochastic-gradient-descent>




Supervised learning



# Forward and backward

Model run a forward to get the output

- It's what actually inference do



NIPS2015 DL-Tutorial (Geoff Hinton, Yoshua Bengio, Yann LeCun)

# Forward and backward

Backpropagation to calculate the gradients, maybe different graph path compare to forward

- Backprop for the activities

$$dC / dX_{i-1} = dC / dX_i * dX_i / dX_{i-1}$$

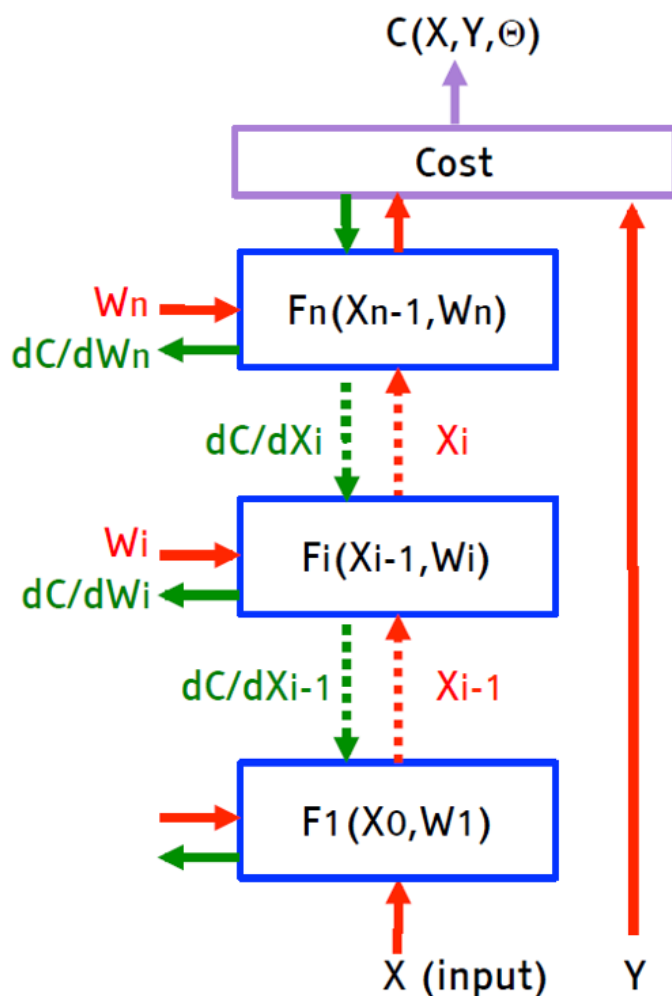
$$dC / dX_{i-1} = dC / dX_i * dF_i(X_{i-1}, W_i) / dX_{i-1}$$

- Backprop for the weights

$$dC / dW_i = dC/dX_i * dX_i / dW_i$$

$$dC / dW_i = dC/dX_i * dF_i(X_{i-1}, W_i) / dW_i$$

NIPS2015 DL-Tutorial (Geoff Hinton, Yoshua Bengio, Yann LeCun)



# Now come back to the code

- Optimizer
  - Define a training process
- Optim Method
  - SGD, Adam...
- Triggers
  - when to validate
  - when to stop training
  - when to checkpoint model
- Batch size
  - Global batch size, should be dividable by the core number
- Validation method
  - Evaluate metric on validation data

# Train a logistic regression model

```
%%time
from bigdl.nn.layer import *
from bigdl.nn.criterion import *
from bigdl.optim.optimizer import *
from bigdl.util.common import *

def logistic_regression(n_input, n_classes):
    # Initialize a sequential container
    model = Sequential()

    model.add(Reshape([28*28]))
    model.add(Linear(n_input, n_classes))
    model.add(LogSoftMax())

    return model
```

# Train a logistic regression model

```
model = logistic_regression(784, 10)
```

```
init_engine()
```

```
optimizer = Optimizer(  
    model=model,  
    training_rdd=train_data,  
    criterion=ClassNLLCriterion(),  
    optim_method="SGD",  
    state={"learningRate": 0.2},  
    end_trigger=MaxEpoch(15),  
    batch_size=2048)
```

```
# Start to train
```

```
trained_model = optimizer.optimize()  
print "Optimization Done."
```

# Inference

```
def map_predict_label(l):  
    return np.array(l).argmax()  
def map_groundtruth_label(l):  
    return l[0] - 1  
  
# Prediction  
predictions = trained_model.predict(test_data)  
imshow(np.column_stack([np.array(s.features).reshape(28,28) for s in  
test_data.take(8)]),cmap='gray'); axis('off')  
print 'Ground Truth labels:'  
print ', '.join(str(map_groundtruth_label(s.label)) for s in test_data.take(8))  
print 'Predicted labels:'  
print ', '.join(str(map_predict_label(s)) for s in predictions.take(8))
```

# Train a CNN model

```
%%time
def build_model(class_num):
    model = Sequential()
    model.add(Reshape([1, 28, 28]))
    model.add(SpatialConvolution(1, 6, 5, 5).set_name('conv1'))
    model.add(Tanh())
    model.add(SpatialMaxPooling(2, 2, 2, 2).set_name('pool1'))
    model.add(Tanh())
    model.add(SpatialConvolution(6, 12, 5, 5).set_name('conv2'))
    model.add(SpatialMaxPooling(2, 2, 2, 2).set_name('pool2'))
    model.add(Reshape([12 * 4 * 4]))
    model.add(Linear(12 * 4 * 4, 100).set_name('fc1'))
    model.add(Tanh())
    model.add(Linear(100, class_num).set_name('score'))
    model.add(LogSoftMax())
    return model
```

# Train a CNN model

```
lenet_model = build_model(10)

import datetime as dt

optimizer = Optimizer(
    model=lenet_model,
    training_rdd=train_data,
    criterion=ClassNLLCriterion(),
    optim_method="SGD",
    state={"learningRate": 0.4, "learningRateDecay": 0.0002},
    end_trigger=MaxEpoch(5),
    batch_size=2048)
```



# Train a CNN model

```
optimizer.set_validation(  
    batch_size=2048,  
    val_rdd=test_data,  
    trigger=EveryEpoch(),  
    val_method=["Top1Accuracy"]  
)  
  
app_name='lenet-'+dt.datetime.now().strftime("%Y%m%d-%H%M%S")  
train_summary = TrainSummary(log_dir='/tmp/bigdl_summaries',  
                             app_name=app_name)  
train_summary.set_summary_trigger("Parameters", SeveralIteration(50))  
val_summary = ValidationSummary(log_dir='/tmp/bigdl_summaries',  
                                app_name=app_name)  
optimizer.set_train_summary(train_summary)  
optimizer.set_val_summary(val_summary)  
print "saving logs to ",app_name
```

# Train a CNN model

```
# Boot training process  
trained_model = optimizer.optimize()  
print "Optimization Done."
```

# Visualize your training

```
loss = np.array(train_summary.read_scalar("Loss"))  
top1 = np.array(val_summary.read_scalar("Top1Accuracy"))
```

```
plt.figure(figsize = (12,12))  
plt.subplot(2,1,1)  
plt.plot(loss[:,0],loss[:,1],label='loss')  
plt.xlim(0,loss.shape[0]+10)  
plt.grid(True)  
plt.title("loss")  
plt.subplot(2,1,2)  
plt.plot(top1[:,0],top1[:,1],label='top1')  
plt.xlim(0,loss.shape[0]+10)  
plt.title("top1 accuracy")  
plt.grid(True)
```

# MODEL TUNING

# Tuning your optimization

- Choose the hyper-parameter of the optimization algorithm carefully
- Hyper-parameter need to adjust when batch size change
- Set log level of `com.intel.analytics.bigdl.optim.DistriOptimizer` to debug to see fine details of your training process
- Use physical core number of your server, which means if hyper-thread is turned on, use half of the v-core number
- Initialize your model correctly
- Use some regularization
- Visualize your training process

# Initialize your model correctly

Model parameter is initialized randomly. You can change how to init them

- Uniform distribution
- Normal distribution
- Constant
- Xavier
- Bilinear

Bad initialization may cause model can't train

# Regularization

Regularization is important to improve model quality

Set it in optimization algorithm

**Python:** `val sgd = new SGD(..., weightDecay = 0.001, ...)`

**Scala:** `sgd = SGD(..., weight_decay = 0.001, ...)`

Set it layer wise

```
Linear(inputN, outputN,  
      wRegularizer = L2Regularizer(0.1),  
      bRegularizer = L2Regularizer(0.1))
```

# The challenge to train deep model

## Gradient vanishing / exploding

- ReLU
- Initialize model correctly (Xavier/pre-trained model)
- Batchnormalization

## Overfitting

- More data (data augmentation)
- Regularization
- Dropout




# Visualize your training process

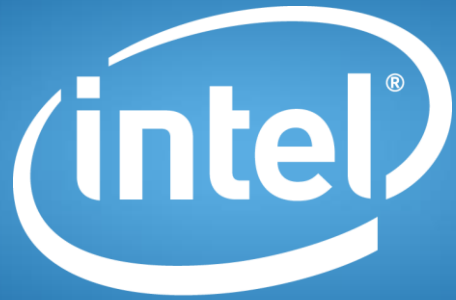
## Turn on persist training summary, Scala

```
val optimizer = Optimizer(...)
...
val logdir = "mylogdir"
val appName = "myapp"
val trainSummary = TrainSummary(logdir, appName)
trainSummary.setSummaryTrigger("Parameters", Trigger.severalIteration(20))
val validationSummary = ValidationSummary(logdir, appName)
optimizer.setTrainSummary(trainSummary)
optimizer.setValidationSummary(validationSummary)
...
val trained_model = optimizer.optimize()
```

```
pip install tensorboard==1.0.0a4
tensorboard --logdir=/tmp/bigdl_summaries
```

# Visualize your training process





Software

# Legal Disclaimer

INFORMATION IN THIS DOCUMENT IS PROVIDED IN CONNECTION WITH INTEL PRODUCTS. NO LICENSE, EXPRESS OR IMPLIED, BY ESTOPPEL OR OTHERWISE, TO ANY INTELLECTUAL PROPERTY RIGHTS IS GRANTED BY THIS DOCUMENT. EXCEPT AS PROVIDED IN INTEL'S TERMS AND CONDITIONS OF SALE FOR SUCH PRODUCTS, INTEL ASSUMES NO LIABILITY WHATSOEVER AND INTEL DISCLAIMS ANY EXPRESS OR IMPLIED WARRANTY, RELATING TO SALE AND/OR USE OF INTEL PRODUCTS INCLUDING LIABILITY OR WARRANTIES RELATING TO FITNESS FOR A PARTICULAR PURPOSE, MERCHANTABILITY, OR INFRINGEMENT OF ANY PATENT, COPYRIGHT OR OTHER INTELLECTUAL PROPERTY RIGHT.

A "Mission Critical Application" is any application in which failure of the Intel Product could result, directly or indirectly, in personal injury or death. SHOULD YOU PURCHASE OR USE INTEL'S PRODUCTS FOR ANY SUCH MISSION CRITICAL APPLICATION, YOU SHALL INDEMNIFY AND HOLD INTEL AND ITS SUBSIDIARIES, SUBCONTRACTORS AND AFFILIATES, AND THE DIRECTORS, OFFICERS, AND EMPLOYEES OF EACH, HARMLESS AGAINST ALL CLAIMS COSTS, DAMAGES, AND EXPENSES AND REASONABLE ATTORNEYS' FEES ARISING OUT OF, DIRECTLY OR INDIRECTLY, ANY CLAIM OF PRODUCT LIABILITY, PERSONAL INJURY, OR DEATH ARISING IN ANY WAY OUT OF SUCH MISSION CRITICAL APPLICATION, WHETHER OR NOT INTEL OR ITS SUBCONTRACTOR WAS NEGLIGENT IN THE DESIGN, MANUFACTURE, OR WARNING OF THE INTEL PRODUCT OR ANY OF ITS PARTS.

Intel may make changes to specifications and product descriptions at any time, without notice. Designers must not rely on the absence or characteristics of any features or instructions marked "reserved" or "undefined". Intel reserves these for future definition and shall have no responsibility whatsoever for conflicts or incompatibilities arising from future changes to them. The information here is subject to change without notice. Do not finalize a design with this information.

The products described in this document may contain design defects or errors known as errata which may cause the product to deviate from published specifications. Current characterized errata are available on request.

Contact your local Intel sales office or your distributor to obtain the latest specifications and before placing your product order.

Copies of documents which have an order number and are referenced in this document, or other Intel literature, may be obtained by calling 1-800-548-4725, or go to: <http://www.intel.com/design/literature.htm>

Intel, Quark, VTune, Xeon, Cilk, Atom, Look Inside and the Intel logo are trademarks of Intel Corporation in the United States and other countries.

\*Other names and brands may be claimed as the property of others.

Copyright ©2015 Intel Corporation.

# Risk Factors

The above statements and any others in this document that refer to plans and expectations for the first quarter, the year and the future are forward-looking statements that involve a number of risks and uncertainties. Words such as “anticipates,” “expects,” “intends,” “plans,” “believes,” “seeks,” “estimates,” “may,” “will,” “should” and their variations identify forward-looking statements. Statements that refer to or are based on projections, uncertain events or assumptions also identify forward-looking statements. Many factors could affect Intel’s actual results, and variances from Intel’s current expectations regarding such factors could cause actual results to differ materially from those expressed in these forward-looking statements. Intel presently considers the following to be the important factors that could cause actual results to differ materially from the company’s expectations. Demand could be different from Intel’s expectations due to factors including changes in business and economic conditions; customer acceptance of Intel’s and competitors’ products; supply constraints and other disruptions affecting customers; changes in customer order patterns including order cancellations; and changes in the level of inventory at customers. Uncertainty in global economic and financial conditions poses a risk that consumers and businesses may defer purchases in response to negative financial events, which could negatively affect product demand and other related matters. Intel operates in intensely competitive industries that are characterized by a high percentage of costs that are fixed or difficult to reduce in the short term and product demand that is highly variable and difficult to forecast. Revenue and the gross margin percentage are affected by the timing of Intel product introductions and the demand for and market acceptance of Intel’s products; actions taken by Intel’s competitors, including product offerings and introductions, marketing programs and pricing pressures and Intel’s response to such actions; and Intel’s ability to respond quickly to technological developments and to incorporate new features into its products. The gross margin percentage could vary significantly from expectations based on capacity utilization; variations in inventory valuation, including variations related to the timing of qualifying products for sale; changes in revenue levels; segment product mix; the timing and execution of the manufacturing ramp and associated costs; start-up costs; excess or obsolete inventory; changes in unit costs; defects or disruptions in the supply of materials or resources; product manufacturing quality/yields; and impairments of long-lived assets, including manufacturing, assembly/test and intangible assets. Intel’s results could be affected by adverse economic, social, political and physical/infrastructure conditions in countries where Intel, its customers or its suppliers operate, including military conflict and other security risks, natural disasters, infrastructure disruptions, health concerns and fluctuations in currency exchange rates. Expenses, particularly certain marketing and compensation expenses, as well as restructuring and asset impairment charges, vary depending on the level of demand for Intel’s products and the level of revenue and profits. Intel’s results could be affected by the timing of closing of acquisitions and divestitures. Intel’s results could be affected by adverse effects associated with product defects and errata (deviations from published specifications), and by litigation or regulatory matters involving intellectual property, stockholder, consumer, antitrust, disclosure and other issues, such as the litigation and regulatory matters described in Intel’s SEC reports. An unfavorable ruling could include monetary damages or an injunction prohibiting Intel from manufacturing or selling one or more products, precluding particular business practices, impacting Intel’s ability to design its products, or requiring other remedies such as compulsory licensing of intellectual property. A detailed discussion of these and other factors that could affect Intel’s results is included in Intel’s SEC filings, including the company’s most recent reports on Form 10-Q, Form 10-K and earnings release.